

# Demystifying Deep Learning in Networking

**Ying Zheng**, Ziyu Liu, Xinyu You, Yuedong Xu      Junchen Jiang



復旦大學



THE UNIVERSITY OF  
CHICAGO

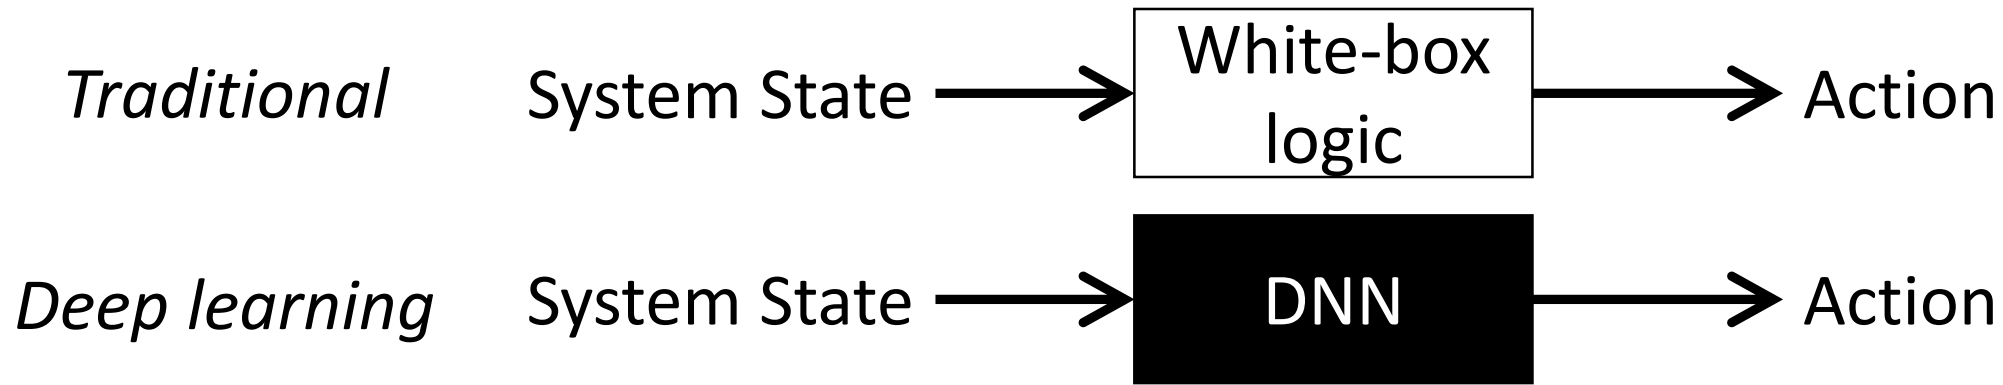
# Why Deep Neural Nets in Networking?

Application	Example gains
Cloud job scheduling	<b>32.4%</b> less job slowdown
Adaptive-bitrate streaming	<b>12-25%</b> better QoE
Routing	<b>23-50%</b> less congestion
Cellular traffic scheduling	<b>14.7%</b> higher utilization


*Many more papers are coming!*

Great promises of DNNs: Two-digit improvements in multiple applications

# So what's wrong with DNNs in networking?



*DNN is a “black box”*



## Problems

- **Hard to confer causality**
- **Hard to trust**
- **Agnostic to domain knowledge**

# Black-Box Problem #1: Hard to confer causality

Sequence of incoming  
jobs (length)

Picked job  
(length)

5, 1, 3, 10, ...

DNN

1

9, 2, 7, 15, ...

DNN

2

*So probably DNN is looking for shortest job*

*But ...*


20, 11, 13, 14

DNN

Wait  
(not 11!)

Correlation  $\neq$  Causality

## Problem #2: Hard to trust





System engineers may favor certainty over performance

# Problem #3: Agnostic to domain-specific knowledge

**Empirical requirement:**

Resource utilization should be below 80%



Hard to imbue a domain-specific rule in a black box.

# Roadmap to white-box DNNs

- How is a decision made?
- When will it fail?
- Can domain-specific knowledge be integrated?

# Outline


- Why try to interpret deep learning in networking
- **An attempt to explain DNN in networking: A case study**
- Examples of worst-case performance of DNN
- Improving performance by domain-specific knowledge



# First attempt: use saliency map to understand DNNs

- Calculate gradient vector


$$w_i = \frac{\partial y_i}{\partial x}$$



- $x$ : input vector
- $y$ : output vector

the  $j^{th}$  element of  $w_i$  shows how much a small change on the  $j^{th}$  feature of  $x$  will change how likely  $i^{th}$  action is picked.

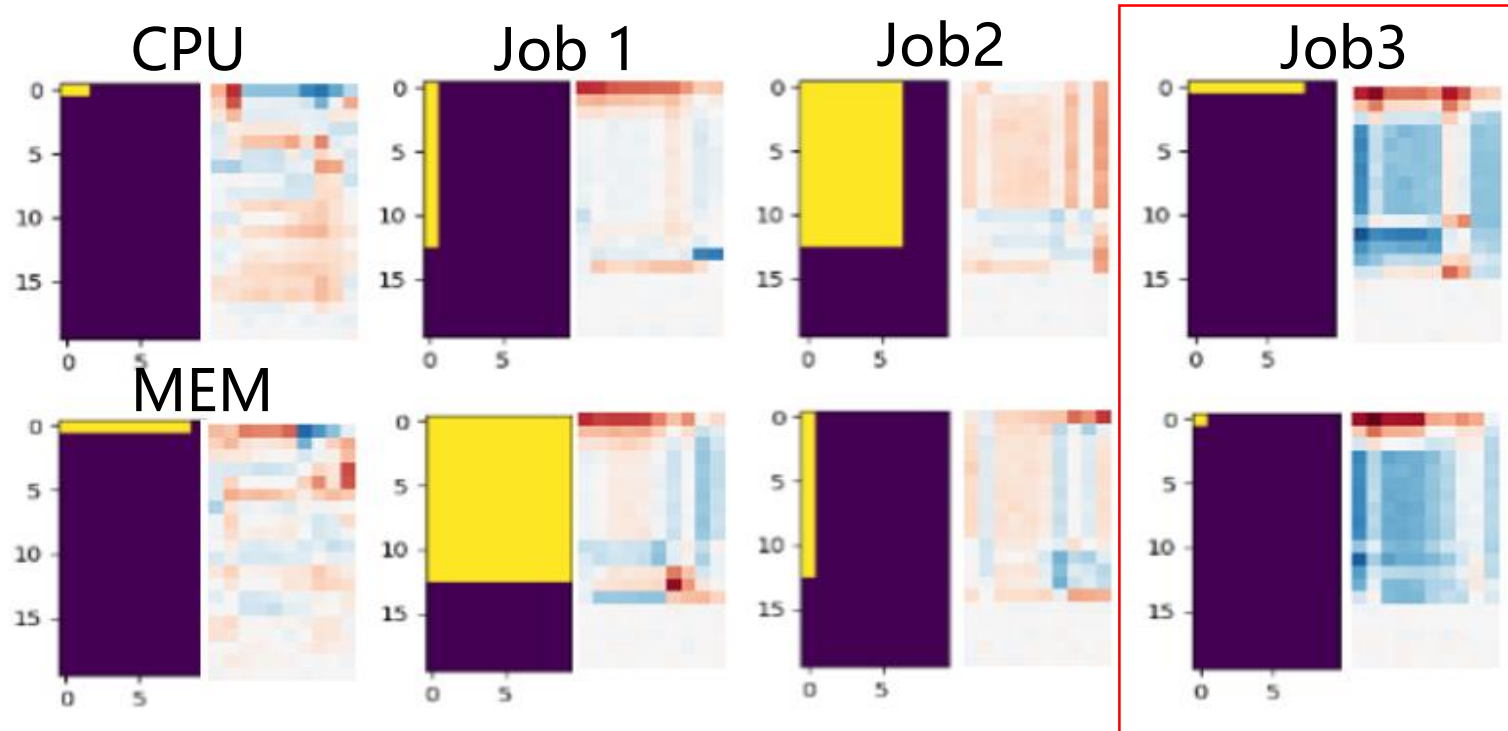
# DeepRM Input/Output: A Case Study



- Input: Remaining resource, Resource requirements
- Output: Act probability over possible actions

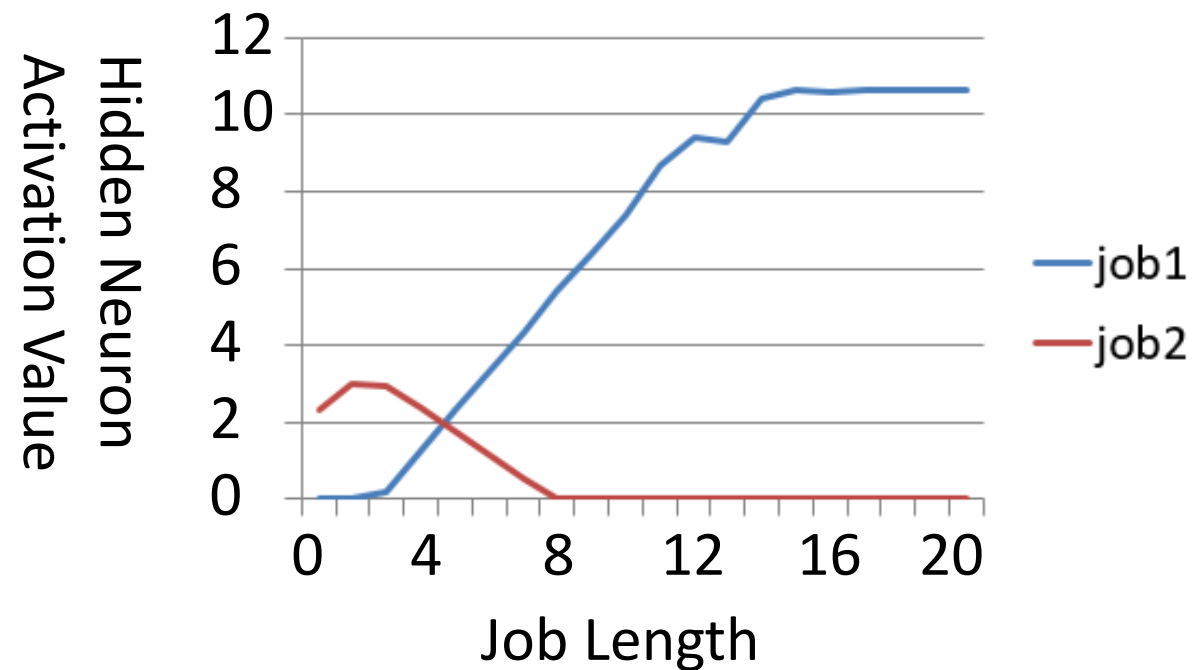
# What features make DeepRM a decision?

- Saliency map



- CPU/RAM occupancy has small impact on the output
- Request time of jobs(outstanding district)

# Relationship between intermediate output and decision



- High-level features have connections to intermediate layer and affect final output

# Outline

- Why try to interpret deep learning in networking
- Reverse-engineer DNN in networking: A case study
- **Examples of worst-case performance of DNN**
- Improving performance by domain-specific knowledge

# When DNN fails in a different workload than what it's trained on

Comparison objective	Testing distribution	DNN	Rule-based algorithm
Average job slowdown	Training distribution	3.73	4.51
		3.25	4.87
		3.93	5.37
		3.03	4.71
	Other distribution	10.57	9.69
		11.44	10.23
		12.02	10.75
		11.48	10.22

When DNN fails *\*even\** in the same workload  
to what it's trained on

Comparison objective	Test data	DeepRM	Rule-based algorithm
Average job slowdown	Normal sequence	1.84	2.01
	Adversarial sequence	1.80	1.46

# Outline

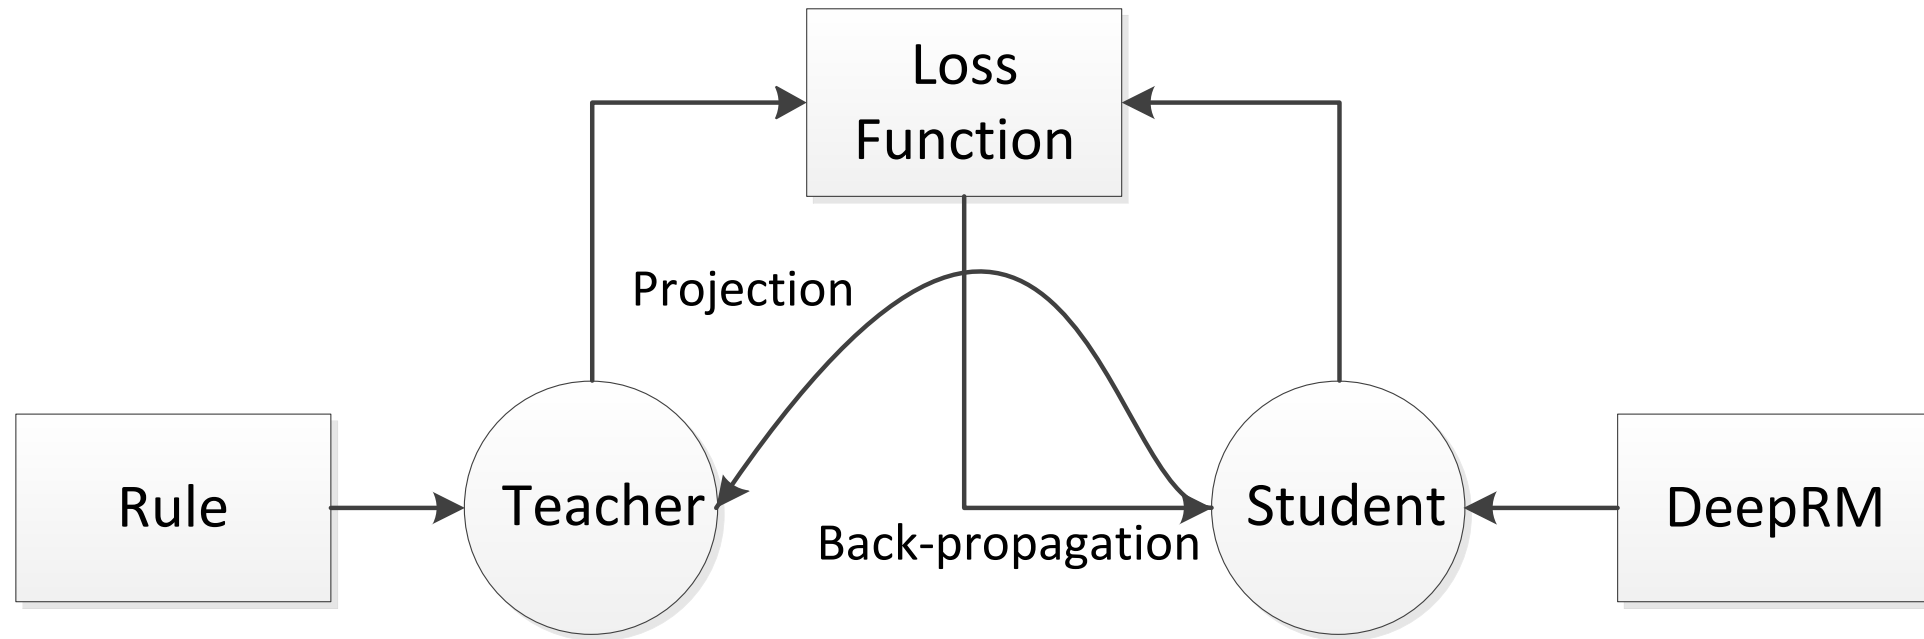
- Why try to interpret deep learning in networking
- Reverse-engineer DNN in networking: A case study
- Examples of worst-case performance of DNN
- Improving performance by domain-specific knowledge



# Domain-specific knowledge might provide a cure

- Better transferability
- Better robustness
- Better training efficiency

# A standard way to incorporate domain knowledge



# Some takeaways, a lot more need to be done

- DNNs in networking achieve superior performance
- DNNs' limitations compel us to try to at least understand it.
- Our preliminary findings:
  - a) One can explain to some extent how networking DNNs work
  - b) DNNs may fail: fits only distribution, vulnerable to noises

Thanks for your listening!